

Spectrum Preserving QUBO Embeddings on a Quantum Annealer

Michael L Rogers
Los Alamos National Laboratory
Los Alamos, New Mexico 87545, USA

(Dated: Aug 21 2019)

Abstract

We show how to embed a QUBO problem on the D-Wave architecture chaining qubits together in a way that avoids numerically perturbing the energy spectrum by introducing a counter-term in the weights for the chained qubits.

Contents

I. Introduction	3
II. Writing a QUBO problem as a “Logical” Hamiltonian	4
III. Embedding the QUBO Problem on the Chimera Chip	7
A. General Considerations	7
B. Chaining Strategies	9
1. The Standard Approach	9
2. The Spectrally Invariant Approach	12
3. Example	13

I. INTRODUCTION

Embedding a QUBO Hamiltonian onto a the D-Wave architecture generally requires coupling or “chaining” several “physical” qubits together to form single “logical” qubits. Chaining is accomplished by adding a penalty term to the coupling interaction of the chained qubits which is meant to result in a lower ground state energy for all of the qubits used in the problem when the qubits in each chain are all correlated (i.e, have the same value) in the ground state. Thus, when the ground state of the total Hamiltonian, consisting of the ground state vector of all the physical qubits, corresponds to the ground state for the Hamiltonian in only the logical qubits, the solution of the embedded problem will correspond to the intended QUBO solution. However, putting the penalty term only into the coupling terms perturbs the total energy in such a way that the ground state with *all* of the physical qubits used in the problem will generally have a different value than the energy with only the logical qubits. This will generally be true even when the chained qubits are correctly correlated so that ground state of the total Hamiltonian corresponds to the desired ground state of the Hamiltonian in the logical qubits only.

This perturbation is highly undesirable because it introduces an effective error into the total energy (relative to the solution energy in logical qubits) which adds cumulatively from each chain and which is dependent on each chains size, and, therefore, upon the particular embedding employed. Furthermore, it is necessary to set the penalty term high enough correlate the qubits in a chain, however, a sufficiently high chaining penalty can increase the chain size dependent energy contribution to a degree that the combined chaining energy terms may dwarf the allowed dynamic range on the parameters. This can cause the values of some of the original strengths and weights desired for the logical qubits to be at or below the level of “noise” on the D-Wave machine. Then the embedded ground state will not correspond to the desired problem solution. And, in many cases, if the chaining penalty is set lower than this, the cumulative error from the different chain sizes in the embedding will add up so that the lowest total energy from some state having one or more “broken chains” is actually *lower* than that corresponding to the desired ground state in logical qubits. Thus, states with broken chains may end up having lower energies than the desired ground state corresponding to the original QUBO problem. In fact, there may be very many states of with broken chains which have a lower total energy than the desired ground state corresponding to QUBO solution. This occurs commonly for the more difficult QUBO problems to solve on the D-Wave, such as those involving the solution of ill-conditioned linear systems, or more computationally complex problems, so it can become impossible to practically solve such problems on a D-Wave machine.

However, it is actually unnecessary to perturb the ground state energy from that of the Hamiltonian in logical qubits to successfully embed the QUBO problem into a physical Hamiltonian with chained qubits. Indeed, it is easy to embed any logical Hamiltonian, without any perturbation at all of its entire spectrum by simply adding a counter-term to the weights for each chained physical qubit, provided one knows the lengths of each chain in the embedding. Then, the ground state of the embedded problem with no broken chains will correspond exactly to the ground state of original QUBO Hamiltonian. The spectrum of the embedded Hamiltonian will then be independent of the value of the chaining penalty, so that it may be set as high as necessary to successfully chain all of the qubits in each logical qubit together.

This counter-term weight is derived and demonstrated in the following section.

II. WRITING A QUBO PROBLEM AS A “LOGICAL” HAMILTONIAN

We assume that some computation problem one wishes to solve has been written as an equivalent Quadratic Binary Unconstrained Optimization (QUBO) problem. The first step in mapping a general QUBO problem onto the D-Wave machine begins with constructing a Hamiltonian that encodes the logical problem in terms of a set of qubits. Next, it will be necessary to “embed” the problem on the chip, first by mapping each logical qubit to a collection or “chain” of physical qubits, and then by determining parameter settings for all the physical qubits, including the chain couplings. By a “logical” Hamiltonian, we mean the original problem Hamiltonian derived for idealized “logical” qubits, which may have arbitrary connectivity. Assume one has already expressed some problem she wishes to solve in QUBO form, i.e., as a quadratic objective function in binary variables whose minimum valued state is the solution of the problem. For the D-Wave machine, this objective function will correspond to the Hamiltonian one applies to the D-Wave machine, and the ground state will be the binary-encoded solution state.

It is useful to review the basic formalism and to establish some notation, which we do in terms of “logical” qubits, assuming arbitrary connectivity. The convention we adopt is that the i th “logical” qubit is denoted with a in upper-case as Q_i , and the r -th “physical” qubit is denoted in lower-case as q_r . Indices i, j, k will be used for logical qubits, and indices r, s, t for physical qubits.

For general QUBO problem with arbitrary connectivity, the connections between the logical qubits can be represented by a graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$, where \mathcal{V} is the vertex set and \mathcal{E} is the edge set. The QUBO *Hamiltonian* on \mathcal{G} is defined by

$$H_{\mathcal{G}}[Q] = \sum_{r \in \mathcal{V}} A_r Q_r + \sum_{rs \in \mathcal{E}} B_{rs} Q_r Q_s , \quad (2.1)$$

with $Q_r \in \{0, 1\}$ for all $r \in \mathcal{V}$. The coefficient A_r is called the *weight* at vertex r , while the coefficient B_{rs} is called the *strength* between vertices r and s . It might be better to call (2.1) the *objective function* rather than the Hamiltonian, as H_G is a real-valued function and not an operator on a Hilbert-space. However, it is easy to map (2.1) in an equivalent Hilbert space form,

$$\hat{H}_G = \sum_{r \in \mathcal{V}} A_r \hat{Q}_r + \sum_{rs \in \mathcal{E}} B_{rs} \hat{Q}_r \hat{Q}_s , \quad (2.2)$$

where $\hat{Q}_r|Q\rangle = Q_r|Q\rangle$ for all $r \in \mathcal{V}$, and $|Q\rangle \in \mathcal{H}$ for Hilbert space \mathcal{H} . The hat denotes an operator on the Hilbert space, and Q_r is the corresponding Eigenvalue of \hat{Q}_r with Eigenstate $|Q\rangle$. Consequently, we can write

$$\hat{H}_G|Q\rangle = H_G[Q]|Q\rangle , \quad (2.3)$$

and we use the terms *Hamiltonian* and *objective function* interchangeably. By the *QUBO problem*, we mean the problem of finding the lowest energy state $|Q\rangle$ of the Hamiltonian (2.2), which corresponds to minimizing Eq. (2.1) with respect to the Q_r . This is, in general, an NP-hard problem uniquely suited to quantum annealing. Rather than sampling all $2^{\#\mathcal{V}}$ possible states, quantum tunneling finds the *most likely* path to the ground state by minimizing the Euclidian action. In the case of the D-Wave 2X chip, the number of distinct quantum states is of order the very large number 2^{1000} , and the ground state is selected from this jungle of quantum states by tunneling to those states with a smaller Euclidean action.

Consider a QUBO problem with R bits of resolution which may have arbitrary connectivity between qubits. Then, the graph for us a problem \mathcal{G} is just the fully connected graph K_R . In terms of vertex and edge sets, we write $K_R = (\mathcal{V}_R, \mathcal{E}_R)$, and Fig. 1 illustrates K_8 and K_4 . The left panel shows the completely connected graph K_8 , with vertex and edge sets

$$\mathcal{V}_8 = \{0, 1, 2, \dots, 7\} \quad (2.4)$$

$$\mathcal{E}_8 = \{\{0, 1\}, \{0, 2\}, \dots, \{0, 7\}, \{1, 2\}, \dots, \{1, 7\}, \dots, \{6, 7\}\} , \quad (2.5)$$

while the right panel shows the K_4 graph,

$$\mathcal{V}_4 = \{0, 1, 2, 3\} \quad (2.6)$$

$$\mathcal{E}_4 = \{\{0, 1\}, \{0, 2\}, \{0, 3\}, \{1, 2\}, \{1, 3\}, \{2, 3\}\} . \quad (2.7)$$

Rather than summing over an ordered edge set, i.e.,

$$H[Q] = \sum_{r \in \mathcal{V}_R} A_r Q_r + \sum_{rs \in \mathcal{E}_R} B_{rs} Q_r Q_s \quad (2.8)$$

$$= \sum_{r=0}^{R-1} A_r Q_r + \sum_{r=0}^{R-1} \sum_{s>r}^{R-1} B_{rs} Q_r Q_s , \quad (2.9)$$




FIG. 1: The left panel shows the fully connected graph K_8 and the right panel shows the corresponding graph K_4 . To perform a calculation to 8-bit accuracy requires the connectivity of K_8 . We take the vertex and edge sets for K_8 to be $\mathcal{V}_8 = \{0, 1, 2, \dots, 7\}$ and $\mathcal{E}_8 = \{\{0, 1\}, \{0, 2\}, \dots, \{0, 7\}, \{1, 2\}, \{1, 3\}, \dots, \{6, 7\}\}$. To perform a calculation to 4-bit accuracy requires K_4 connectivity, and similarly, the vertex and edge sets for K_4 are $\mathcal{V}_4 = \{0, 1, 2, 3\}$ and $\mathcal{E}_4 = \{\{0, 1\}, \{0, 2\}, \{0, 3\}, \{1, 2\}, \{1, 3\}, \{2, 3\}\}$.

we find it convenient to sum over all values of r and s taking B_{rs} to be symmetric. In this case, the double sum differs by a factor of two relative to summing over the edge set of the graph,

$$H[Q] = \sum_{r=0}^{R-1} A_r Q_r + \sum_{r=0}^{R-1} \sum_{s=0}^{R-1} \frac{1}{2} B_{rs} Q_r Q_s . \quad (2.10)$$

Furthermore, for $r = s$, there will be a linear contribution from the idempotency condition $Q_r^2 = Q_r$, so that

$$H[Q] = \sum_{r=0}^{R-1} \left[A_r + \frac{1}{2} B_{rr} \right] Q_r + \sum_{r=0}^{R-1} \sum_{s \neq r, s=0}^{R-1} \frac{1}{2} B_{rs} Q_r Q_s . \quad (2.11)$$

We can write this as

$$H[Q] = \sum_{r=0}^{R-1} \tilde{A}_r Q_r + \sum_{r=0}^{R-1} \sum_{s \neq r, s=0}^{R-1} \tilde{B}_{rs} Q_r Q_s . \quad (2.12)$$

III. EMBEDDING THE QUBO PROBLEM ON THE CHIMERA CHIP

A. General Considerations

The D-Wave Chimera chip consists of coupled bilayers of micro rf-SQUIDs overlaid in such a way that, while relatively easy to fabricate, results in a fairly limited set of physical connections between the qubits. However, by *chaining* together well chosen qubits in a positively correlated manner, this limitation can largely be overcome. The process of chaining requires that we (i) embed the logical graph onto the physical graph of the chip (for example K_4 onto C_8) and that we (ii) assign weights and strengths to the physical graph embedding in such a way as to preserve the ground state of the logical system. These steps are called graph embedding and Hamiltonian embedding, respectively.




FIG. 2: The left panel illustrates the bipartate graph C_8 in *column* format, while the right panel illustrates the corresponding graph in *cross* format, often called a Chimera graph. The gray lines represent direct connections between qubits. The cross format is useful since it minimizes the number intersecting connections. The use of red and blue dots emphasize the bipartate nature of C_8 , as every red dot is connected to every blue dot, while none of the red and blue dots are connected to one another. The vertex set of C_8 is taken to be $\mathcal{V}_8 = \{1, 2, \dots, 8\}$ and edge set is $\mathcal{B}_8 = \{\{1, 5\}, \{1, 6\}, \{1, 7\}, \{1, 8\}, \{2, 5\}, \{2, 6\} \dots \{7, 8\}\}$.

Let us explore the connectivity of the D-Wave Chimera chip in more detail. The D-Wave architecture employs the C_8 bipartate Chimera graph as its most basic unit of connectivity. This *unit cell* is illustrated in Fig. 2, and consists of 8 qubits connected in a 4×4 bipartate manner. The left panel of the figure uses a *column* format in laying out the qubits, and the right panel illustrates the corresponding qubits in a *cross* format, where the gray lines represent the direct connections between the qubits. The cross format is useful since it minimizes the number intersecting connections. The complete two dimensional chip is produced by replicating C_8 along the vertical and horizontal directions, as illustrated in Fig. 3, thereby




FIG. 3: The left panel shows the connectivity between four C_8 bipartite Chimera zones, and the right panel illustrates how multiple C_8 graphs are stitched together along the vertical and horizontal directions to provide thousands of possible qubits. A limitation of this connectivity strategy is that red and blue zones cannot communicate directly with one another, as indicated by the black crossed arrows. The purpose of *chaining* is to allow communication between the read and blue qubits.

providing a chip with thousands of qubits. The connections between qubits are limited in two ways: (i) by the connectivity of the basic unit cell C_8 and (ii) by the connectivity between the unit cells across the chip. The bipartite graph $C_8 = (\mathcal{V}_8, \mathcal{B}_8)$ is formally defined by the vertex set $\mathcal{V}_8 = \{1, 2, \dots, 8\}$, and the edge set

$$\begin{aligned} \mathcal{B}_8 = \{ & \{1, 5\}, \{1, 6\}, \{1, 7\}, \{1, 8\}, \{2, 5\}, \{2, 6\}, \{2, 7\}, \{2, 8\}, \\ & \{3, 5\}, \{3, 6\}, \{3, 7\}, \{3, 8\}, \{4, 5\}, \{4, 6\}, \{4, 7\}, \{4, 8\} \}. \end{aligned} \quad (3.1)$$

The set \mathcal{B}_8 represents the connections between a given red qubit and the corresponding blue qubits in the Figures. The red and blue dots illustrate the bipartite nature of C_8 , as every red dot is connected to every blue dot, while none of the blue and red dots are connected to one another.

We will denote the *physical* qubits on the D-Wave chip by q_ℓ . For the D-Wave 2000Q there is a maximum of 2048 qubits, while the D-Wave 2X has 1152 qubits. For the example calculation in this text, we only use 10 to 50 qubits. The physical Hamiltonian or objective function takes the form

$$H[q] = \sum_{\ell} a_{\ell} q_{\ell} + \sum_{\ell \neq m} 2b_{\ell m} q_{\ell} q_m, \quad (3.2)$$

where we have introduced a factor of 2 in the strength to account for the symmetric summation over r and s . We will call the qubits Q_r of the previous section the *logical qubits*. To write a program for the D-Wave means finding an embedding of the logical problem onto

the physical collection of qubits q_ℓ . If the connectivity of the Chimera graphs were large enough, then the logical qubits would coincide exactly with the physical qubits. However, since the graph C_8 possesses less connectivity than K_4 , we must resort to chaining on the D-Wave, even for 4-bit resolution. Figure 4 illustrates the K_4 embedding used by our algorithm, where, as before, the left panel illustrates the bipartate graph in column format, and the right panel illustrates the corresponding graph in cross format.



FIG. 4: The K_4 embedding onto C_8 used in our implementation of 4-bit of division on the D-Wave. The blue lines represent normal connections between qubits, while the red double-lines represent chained qubits, that is to say, qubits that are strictly correlated (and can thereby represent a single logical qubit at a higher level of abstraction). The qubits 1-6 are chained together, as are the qubits 3-8.

B. Chaining Strategies

1. The Standard Approach

In Fig. 4 we have labeled the physical qubits by $\ell = 1, 2, 3 \dots 8$, and we wish to map the logical problem involving $Q_r Q_s Q_t$ onto the four physical qubits $q_5 q_1 q_6 q_2$. The embedding requires that we *chain* together the two qubits 1-6 and 3-8, respectively. We may omit qubits 4 and 7 entirely. As illustrated in Fig. 5, the physical qubits q_1 and q_6 are *chained* together to simulate a single logical qubit Q_t , while qubits q_5 and q_2 are mapped directly to the logical qubits Q_r and Q_s , respectively. Qubit q_5 is assigned the weight $a_5 = A_r$ and the coupling between q_5 and q_1 is assigned the value $b_{51} = B_{rt}$. Similarly for qubit q_2 , the vertex is assigned weight $a_2 = A_s$, and strength between q_2 and q_6 is $b_{26} = B_{st}$. We must now distribute the logical qubit Q_t between q_1 and q_6 by assigning the values a_1, a_6 and b_{16} . We distribute the weight A_t uniformly between qubits q_1 and q_2 , giving $a_1 = A_t/2$ and




FIG. 5: The left panel shows three logical qubits Q_r , Q_s , Q_t with connectivity between r - t and t - s . The box surrounding qubit t means that it will be modeled by a linear chain of physical qubits, as illustrated in the right panel. The labeling is taken from Fig. 4 for qubits 5-1-6-2, where Q_r is mapped to q_5 , Q_s is mapped to q_2 , and Q_t is split between q_1 and q_6 . Qubits q_1 and q_6 are chained together to simulate the single logical qubit Q_t , while qubits Q_r and Q_s map directly onto physical qubits q_5 and q_2 .

$a_6 = A_t/2$. We must now choose b_{16} . In this standard approach, as described in the D-Wave documentation, this is accomplished, this is simply the negative chain penalty term, $-\alpha$.

Using superscripts to denote *logical qubit* indices, we may summarize the standard SAPI approach as,

$$a_r^t = A_t/N_t \quad (3.3)$$

$$b_{i,j}^{r,s} = \begin{cases} B_{r,s}/\nu_{r,s}, & \text{if } q_i \in Q_r \text{ and } q_j \in Q_s \\ 0, & \text{otherwise} \end{cases} \quad (3.4)$$

$$b_{i,i+1}^t = \begin{cases} -\alpha, & \text{if } q_i, q_{i+1} \in Q_t \\ 0, & \text{otherwise} \end{cases} \quad (3.5)$$

where,

$$N_t := \text{The number of physical qubits in the logical qubit } Q_t \quad (3.6)$$

$$\nu_{r,s} := \text{The number of physical qubits coupling logical qubits } Q_r \text{ and } Q_s. \quad (3.7)$$

Usually, the default SAPI embedding will give all $\nu_{r,s} = 1$, that is, only one pair of physical qubits is actually coupled to embody the logical qubit coupling, $B_{r,s}$. However, occasionally the SAPI embedding function will couple together two qubits across a pair of logical qubits, in which case it will set the value of each coupling term to the above form with $\nu = 2$. We will assume from now on that all logical qubits are coupled by only one pair of physical qubits, but this is not essential to the following discussion.

Note that this standard chaining scheme will generally give a total energy that is different than that for the corresponding logical Hamiltonian. To see this, first assume that none of the chains are broken, so that $q_i = Q_r$ for all q_i contain within the logical qubit Q_r . Then, the ground state for this chaining scheme will differ from the ground state in terms of logical qubits by an additive perturbation. For a problem with N logical qubits, each consisting of N_t physical qubits that are chained together within logical qubit Q_t , we can see by inspection that the perturbed ground state energy can be written as,

$$E_0[q_i^0] = E_0[Q_t^0] + \Delta E \quad (3.8)$$

where,

$$\Delta E := \sum_{chains} -\alpha \times (\text{the number of links}) \quad (3.9)$$

$$= -\alpha \times \sum_{t=1}^N (N_t - 1) \quad (3.10)$$

where, Here, the subscript 0 refers to values in the ground state. The expression derives from the fact that there is a contribution of $-\alpha$ from each “link” in the chain, and there are $N_t - 1$ links. Note that this relation would actually hold for every state in the spectrum of the Hamiltonian in logical qubits when there are no unbroken chains.

Observe two things: (1) The ground state energy with all the physical qubits included will be lower than that for that for the logical qubits when there are no unbroken chains, and; (2) The magnitude of the perturbation, ΔE , depends on the values of the particular N_t values, which depends on the particular graph embedding employed. So, one problem is that the perturbation is a function of the graph embedding. The same problem with a different embedding onto the Chimera graph could give a different energy perturbation.

But note, also, that, frequently, there may be very small energy differences resulting from reversals of only a few of the physical qubits in ground state. This typically occurs when there are many near cancellations in energy, so that simultaneous reversal of a pair or group of qubit values within possibly different chains may result in a net lower energy if the corresponding chain perturbation is lower than the other terms in the Hamiltonian involving those qubits. This will generally happen when the chaining penalty is too small to compensate for these energy differences. This possibility should be obvious by just considering that the chaining penalty must be set large enough to bias the physical qubits to positive correlations or they will not chain together. However, the chaining penalty cannot be made arbitrarily large, either. If there are many states of the embedded system with energies close

to the ground state, the chaining penalty may need to be set to tens or hundreds of times the values some of the A_t and $B_{r,s}$ parameters to chain physical qubits into logical qubits. But then, when the Hamiltonian is normalized to values within the dynamic range of the machine, the energies from of the some logical qubits couplings may become negligible, i.e., those couplings will no longer contribute significantly to the total energy. In such cases, the solution may not represent the solution of the intended QUBO problem at all, although there may be no broken chains.

For a QUBO problem to be effectively solveable on the D-Wave quantum annealer, there must be feasible range of chaining penalties; large enough to chain the physical qubits together into logical qubits, but not drastically larger than the values of the strengths and weights for the logical qubits. And, for some QUBO problem Hamiltonians, typically those with very high spectral densities near the ground state, this is simply not possible. Some QUBO problems may either require longer annealing times, or more reads, or they may simply not be practically solvable on the annealer. But it is difficult to know when a hard QUBO problem may yet be solved by longer annealing times or more cycles if the chaining perturbation error has not been accounted for, because one generally cannot objectively determine, a priori, whether the reason it's not converging is primarily due to computability limitations, or the physical limitations of the device, such as thermal noise, or simply from the numerical errors resulting from the chaining perturbation. However, the latter is simply an unnecessary numerical error which is caused by trying to set the parameters without knowing anything about the chain sizes due to the particular graph embedding. Yet, this information is readily available after one has done the necessary step of embedding the graph. Therefore, this chaining information should to be used to correct the chaining perturbation if one wants to have the best chance of solving any potentially difficult QUBO problem.

We address how to use the chain sizes to set the parameters in order to keep the spectrum invariant in the following section.

2. *The Spectrally Invariant Approach*

Clearly, to keep the spectrum invariant, we must eliminate the perturbation from the chaining penalty. To do that, we simply shift the values of the weights, a_i , using information that we already must have from embedding the graph onto the Chimera chip. We only need to know the length of each chain (i.e., the size of each logical qubit), which is known from the graph embedding step. Then, it is straight-forward to add a counter-term to the weights for the physical Hamiltonian to subtract the perturbation from the chaining penalty. The sum over link in a the chain for logical qubit Q_t , gives a perturbation of,

$$\Delta E_{Q_t} = -\alpha \times (N_t - 1). \quad (3.11)$$

As in equation 3.3, for each physical qubit, q_i within logical qubit Q_t we may still choose the weights to have a contribution determined by the logical qubit weights, of the form,

$$a_i \propto \frac{A_t}{N_t}, \quad (3.12)$$

up to a constant term chosen to cancel the chain coupling perturbation. Clearly, the simplest counter-term that will accomplish this is just

$$ct_i^t = \alpha \times \frac{(N_t - 1)}{N_t}. \quad (3.13)$$

Therefore, in our spectrum preserving weighting scheme, we simply add this counter-term to the weights in equations 3.3, giving,

$$a_r^t = \frac{A_t}{N_t} + \alpha \times \frac{(N_t - 1)}{N_t}. \quad (3.14)$$

3. Example

We now illustrate the affect of applying the counter-term vs not applying the counter-term when embedding a Hamiltonian. The Hamiltonian in this is example is for solving a 2-dimensional linear equation defined by an ill-conditioned 2x2 matrix. In such cases, the energy spectrum will have a large number of states packed close in energy to the ground state. The energy spectrum is plotted in Fig. 6, below.

Clearly, the correctly for inhomogeneous chain sizes matters a great deal. Note that, with the correction given by equation 3.14, the energy spectrum (summed over the chained qubits) is exactly equivalent to the energy spectrum for the unchained (logical) Hamiltonian. However, without the correction added to the weights, there is a very significant difference. In fact, it changes the shape of the energy spectrum to such a degree that the lowest energy state is significantly different from the ground state energy of the unchained Hamiltonian, so that the uncorrected embeddiing could not be expected to yield the correct ground state for the original Hamiltonian with the embedded Hamiltonian on the D-Wave quantum annealer.




FIG. 6: The energy surface for a certain QUBO Hamiltonian. This is plotted for three cases: (1) The unembedded original Hamiltonian, in terms of “logical” qubits (blue); (2) the embedded Hamiltonian with the corrected weightings, from equation 3.14 (red), and; (3) the embedded Hamiltonian without the corrected weightings (green).